

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
import tensorflow as tf
import datetime,os
from tensorflow import keras
from keras.models import Model
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.moun

```
#!gdown --id 1Z4TyI7FcFVEx8qd14j09qxvxaqLSqoEu
```

```
#!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT
```

```
#get_ipython().system_raw("unrar x rv1-cdip.rar")
```

```
!unrar x "/content/drive/My Drive/rv1-cdip.rar" "/content/drive/My Drive/image_data_TL"
```

```
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/x/yxx05c00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/x/yxx65c00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/x/yxx86e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/x/yxx86e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/y OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/y/yxy13a00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/y/yxy13a00/50
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/y/yxy58e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/y/yxy58e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz48e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz48e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz71f00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz71f00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz79e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz79e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz85c00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz85c00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz95e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/x/z/yxz95e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/a OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/a/yya60d00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/a/yya60d00/50
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb10e00/20
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb10e00/80
```

Saved successfully!

```

Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb44a00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb44a00/9
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb64c00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb64c00/8
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb80c00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb80c00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb92f00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/b/yyb92f00/t
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc55f00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc55f00/0
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc58d00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc58d00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc88e00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/c/yyc88e00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd01e00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd01e00/8
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd02f00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd02f00/o
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd15c00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/d/ydd15c00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/e OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/e/yye52f00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/e/yye52f00/t
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/e/yye67c00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/e/yye67c00/20
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/f OK
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/f/yyf89d00 (
Extracting    /content/drive/My Drive/image_data_TL/data_final/imagesy/y/y/f/yyf89d00/50
Creating      /content/drive/My Drive/image_data_TL/data_final/imagesv/v/v/g OK

```

#now we load all the labels

```
df = pd.read_csv('/content/drive/My Drive/labels_final.csv')
df.head()
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imageso/o/r/p/ropb0000/517511301+-1301.tif	2
3	imageso/o/r/p/ropb0000/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

Saved successfully!

#from above we can see that the labels of the images are in numbers. So, with the info in ref #convert them understandable categories through their names

```
dict_of_labels = {
    0: 'letter',
    1: 'form',

```

```

2: 'email',
3: 'handwritten',
4: 'advertisement',
5: 'scientific report',
6: 'scientific publication',
7: 'specification',
8: 'file folder',
9: 'news article',
10: 'budget',
11: 'invoice',
12: 'presentation',
13: 'questionnaire',
14: 'resume',
15: 'memo'

}

```

```

df['label']=df['label'].apply(lambda x:dict_of_labels[x])
df.head()

```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	handwritten
1	imagesl/l/x/t/lxt19d00/502213303.tif	handwritten
2	imagesx/x/e/d/xed05a00/2075325674.tif	email
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	handwritten
4	imagesq/q/z/k/qzk17e00/2031320195.tif	specification

```

from keras_preprocessing.image import ImageDataGenerator #import the module
imgdatagen = ImageDataGenerator(rescale=1/255,validation_split=0.2)

```

```

import os
os.getcwd()

```

```

'/content'

```

Saved successfully!

```

train_gen = imgdatagen.flow_from_dataframe(dataframe=df,directory='/content/drive/My Drive/im
#directory: string,path to the target directory that contains all the images mapped in th

```

```

Found 9600 non-validated image filenames belonging to 16 classes.

```

```

val_gen = imgdatagen.flow_from_dataframe(dataframe=df,directory='content/drive/My Drive/image

```

```

Found 9600 non-validated image filenames belonging to 16 classes.

```

```

from keras.layers import Input,Dense,Lambda,Flatten #The Lambda layer exists so that arbitra
from keras.models import Model
from keras.applications.vgg16 import VGG16
from keras.applications.vgg16 import preprocess_input
from keras.preprocessing import image
from keras.layers import Dense,MaxPool2D,Conv2D,Flatten
from keras.callbacks import Callback
from keras.callbacks import TensorBoard

```

```
%load_ext tensorboard
```

```

logdir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d=%H%M%S")
tensorboard_callback = TensorBoard(log_dir = logdir,histogram_freq=1)

```

```

image_size=[256,256]
model = VGG16(input_shape=image_size + [3],weights='imagenet',include_top=False) #input_shape
model.summary()

```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0

Saved successfully!



```

block5_conv1 (Conv2D)      (None, 16, 16, 512)      2359808
block5_conv2 (Conv2D)      (None, 16, 16, 512)      2359808
block5_conv3 (Conv2D)      (None, 16, 16, 512)      2359808
block5_pool (MaxPooling2D) (None, 8, 8, 512)        0

```

```

=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0

```

## Model 1

```

for layer in model.layers:
    layer.trainable = False
x= model.output
x=Conv2D(filters=512,kernel_size=(3,3),padding='same',activation='relu')(x)
x=MaxPool2D(2,2)(x)
x=Flatten()(x)
x=Dense(256,activation='relu')(x)
x=Dense(128,activation='relu')(x)
output=Dense(16,activation='softmax')(x)
model_1 = Model(inputs=model.input,outputs=output)
model_1.compile(loss='categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])
model_1.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080

Saved successfully!



block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d (Conv2D)	(None, 8, 8, 512)	2359808
max_pooling2d (MaxPooling2D)	(None, 4, 4, 512)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 256)	2097408
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 16)	2064

```

=====
Total params: 19,206,864
Trainable params: 4,492,176

```

```

train_steps = train_gen.n//train_gen.batch_size
val_steps = val_gen.n//val_gen.batch_size
print(train_steps,val_steps)

```

```
300 300
```

Saved successfully!



steps\_per\_epoch=train\_steps,epochs=3,validation\_data=val\_ge

pass

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning: `Model.fit_
```

```
Epoch 1/3
```

```
300/300 [=====] - ETA: 0s - loss: 1.1655 - accuracy: 0.6402
```

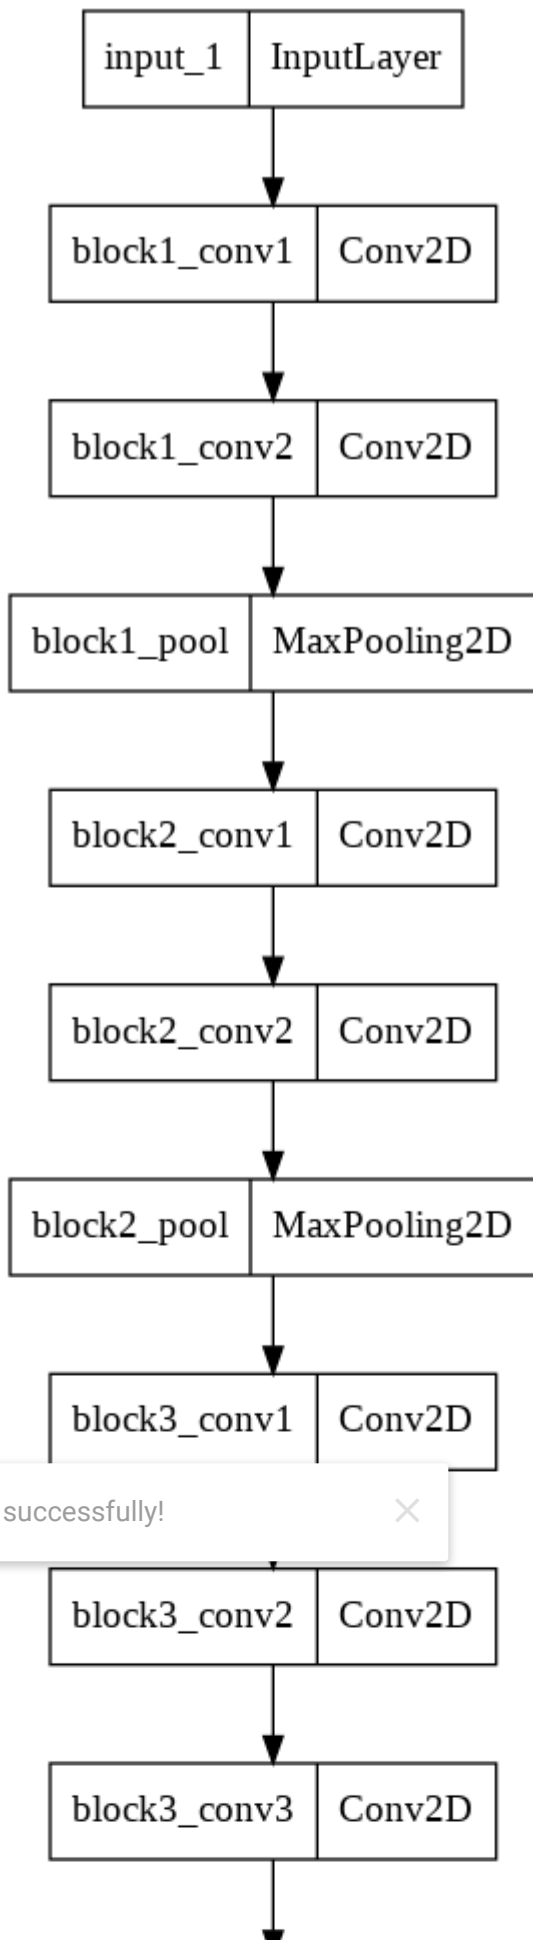


#graph

```
tf.keras.utils.plot_model(model_1,to_file='model_1.png',show_shapes=False,show_layer_names=Tr
```

Saved successfully!






Saved successfully!



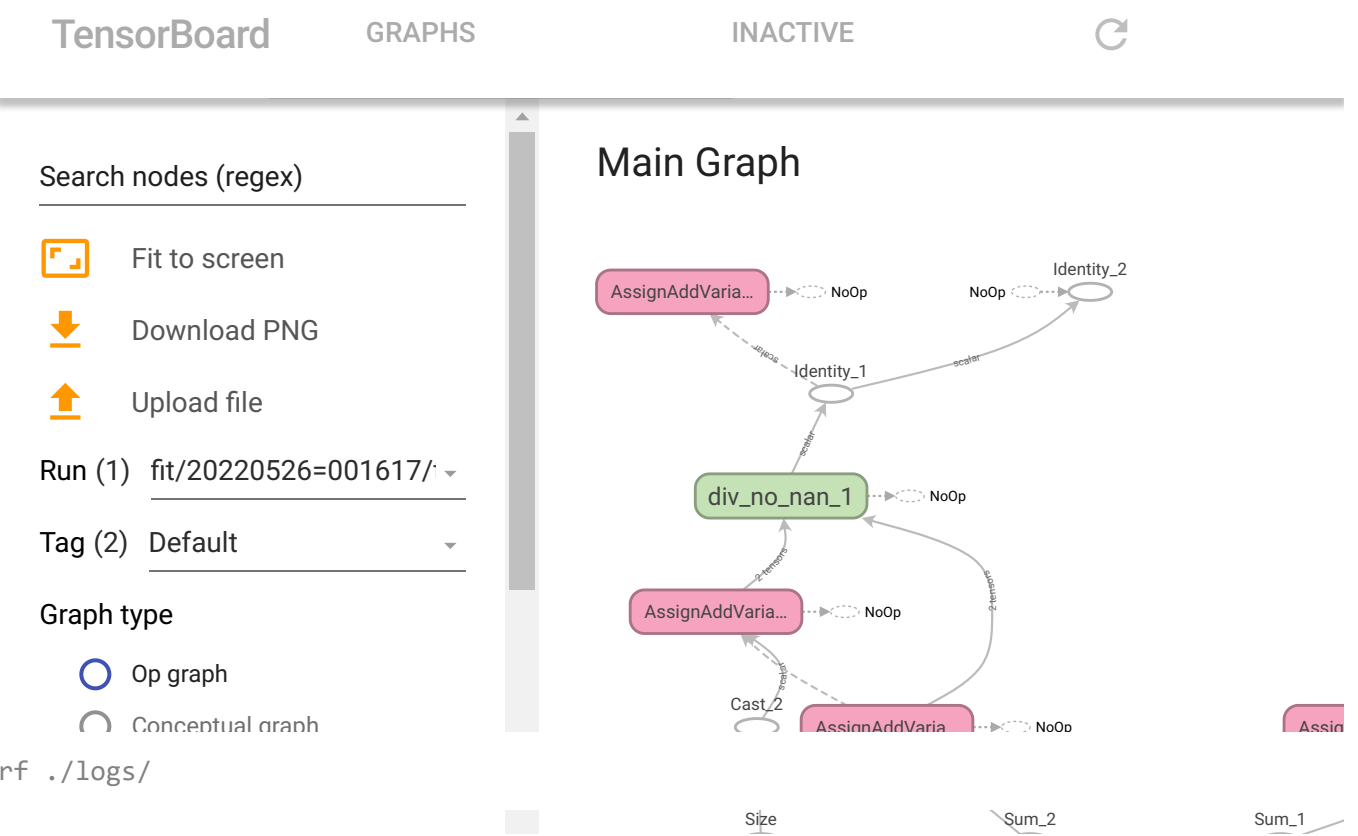




```
%tensorboard --logdir logs
```

Saved successfully!





Key takeaway:

- 1. The total parameters to trainable parameter ratio is around 4.3.
- 2. For that combination, we get an accuracy of 64%

Model 2

```
for layer in model.layers:
    layer.trainable = False
x= model.output
x=Conv2D(filters=4096,kernel_size=8,strides=1,activation='relu')(x)
x=Conv2D(filters=4096,kernel_size=1,strides=1,activation='relu')(x)
model_2=model(inputs=model.input,outputs=output)
model_2.compile(loss='categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])
model_2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792

block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	134221824
conv2d_1 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 16)	65552

Saved successfully!



trainable params: 151,068,688  
Non-trainable params: 14,714,688


try:

```
model_2.fit_generator(train_gen, steps_per_epoch=train_steps, epochs=3, verbose=1, validation_d
except FileNotFoundError:
    pass
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning: `Model.fit_
```

```
Epoch 1/3
```

```
13/300 [>.....] - ETA: 15:53:10 - loss: 1.0707 - accuracy: 0.65
```



```
'''try:
```


```
    model_2.fit_generator(train_gen, steps_per_epoch=train_steps, epochs=3, verbose=1, validation_d  
except FileNotFoundError:
```

```
    pass '''
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning: `Model.fit_
```

```
Epoch 1/3
```

```
300/300 [=====] - ETA: 0s - loss: 2.1016 - accuracy: 0.4876
```

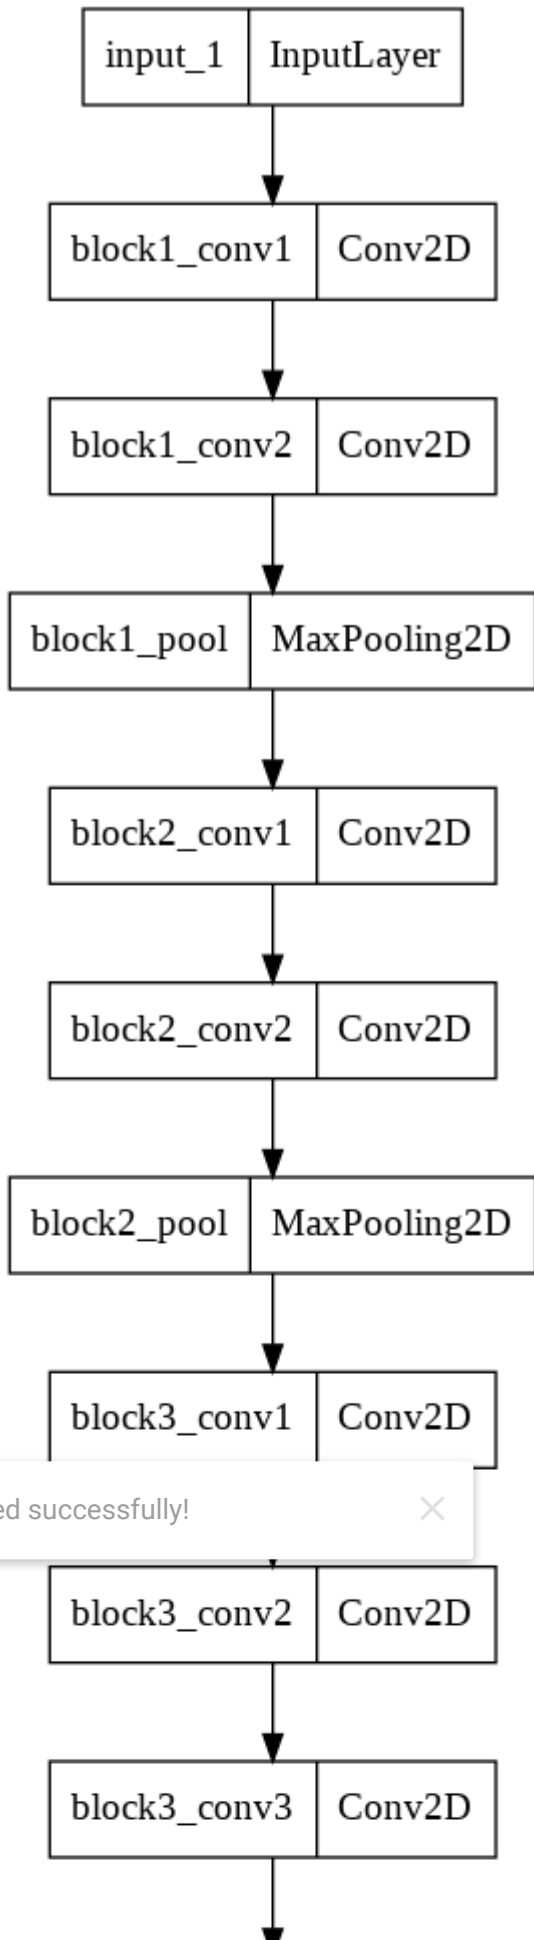


```
#graph
```

```
tf.keras.utils.plot_model(model_2, to_file='model_2.png', show_shapes=False, show_layer_names=Tr
```


Saved successfully!





Saved successfully!



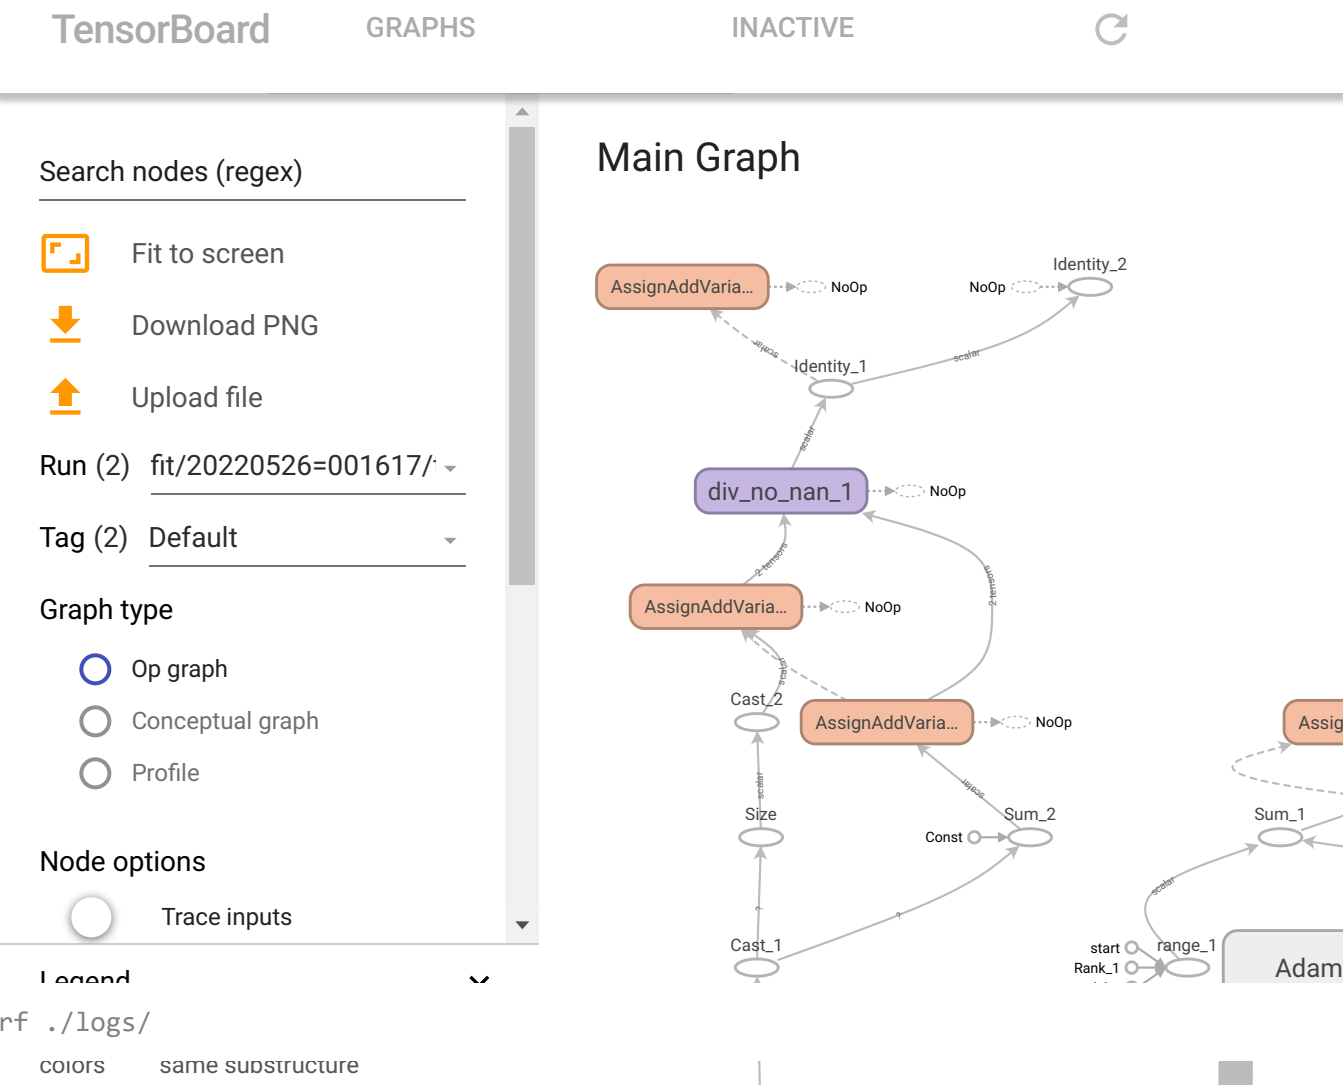


```
%tensorboard --logdir logs
```

Saved successfully!



Reusing TensorBoard on port 6006 (pid 1029), started 1:52:34 ago. (Use '!kill 1029' to kill it.)



Key takeaway:

- 1. The total parameter to trainable parameter ratio is around 1.097 (very close).
- 2. For that ratio, we see the accuracy is around 65.14% (though the 2nd epoch was interrupted abruptly due to RAM insufficiency).
- 3. Here, what we need to take note is that we have paramters in this model than model\_1. It's

Saved successfully!

- Model is secondary to the fact that most parameters are
- 4. So, we can conclude that more trainable parameters have positive effect on the accuracy.

Model 3

```
for layer in model.layers[-6:]:
```

```
layer.trainable = True
```

```
x = model.output
x = Conv2D(filters=4096,kernel_size = 8,strides=1,activation='relu')(x)
x = Conv2D(filters=4096,kernel_size = 1,strides=1,activation='relu')(x)
x = Flatten()(x)
output = Dense(16,activation='softmax')(x)
model_3 = Model(inputs=model.input,outputs=output)
model_3.compile(loss='categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])
```

```
model_3.summary()
```

Model: "model\_3"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0

Saved successfully!





conv2d_6 (Conv2D)	(None, 1, 1, 4096)	134221824
conv2d_7 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten_3 (Flatten)	(None, 4096)	0
dense_3 (Dense)	(None, 16)	65552

```
=====
Total params: 165,783,376
Trainable params: 160,507,920
Non-trainable params: 5,275,456
=====
```

---


try:

```
model_3.fit_generator(train_gen, steps_per_epoch = train_steps, epochs=3, validation_data=val_
except FileNotFoundError:
    pass
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning: `Model.fit_
```

Epoch 1/3

```
300/300 [=====] - ETA: 0s - loss: 2.7727 - accuracy: 0.0625
```



```
tf.keras.utils.plot_model(model_3, to_file='model_3.png', show_shapes=False, show_layer_names=Tr
```

Saved successfully!

